



Advanced Course on Agentic AI: Single-Agent & Multi-Agent Systems

Module 1: Introduction to Agentic AI

- What is Agentic AI? From Traditional AI to Autonomous AI
- Single-Agent vs Multi-Agent Systems (MAS)
- Agentic AI in Large Language Models (LLMs)
- Applications of Agentic AI in Automation, Research, and Business

Hands-on:

- ✓ Running a basic autonomous agent using OpenAI API
 - ✓ Setting up a local agent with Python
-

Module 2: Building Single-Agent Systems

- Single-Agent AI: Concepts & Architectures
- Decision-Making in Single-Agent AI
- Memory & Long-Term Planning
- Integrating LLMs with Agents (LangChain, OpenAI, Hugging Face)
- Handling User Inputs & Actions with Tool-Use Capabilities

Hands-on:

- ✓ Building a goal-driven single AI agent with LangChain
 - ✓ Implementing memory-based reasoning with FAISS
-

Module 3: Multi-Agent Systems (MAS)

- What are Multi-Agent Systems (MAS)?
- Agent Communication & Coordination
- Role Assignment in Multi-Agent Environments
- Task-Oriented AI vs Autonomous Decision-Making
- Swarm Intelligence & Distributed Agents

Hands-on:

- ✓ Creating a multi-agent research assistant
 - ✓ Implementing a task-based workflow with CrewAI
-

Module 4: AI Agents with CrewAI

- Overview of CrewAI: AI Agents Working in Teams
- Role-Based Agent Assignments
- Orchestrating Task Execution Between Agents
- Implementing Workflow Pipelines with CrewAI

Hands-on:

- ✓ Creating an AI-powered content generation team
 - ✓ Using CrewAI for multi-step research tasks
-

Module 5: SmolAgent – Lightweight AI Agents

- Introduction to SmolAgent: Minimalistic AI Agents
- When to Use SmolAgent vs Heavyweight AI Agents

- **Optimizing AI Agents for Cost and Performance**
- **Combining SmolAgent with LLMs for Fast Execution**

Hands-on:

- ✓ Deploying a SmolAgent-based chatbot
 - ✓ Running SmolAgent on an edge device
-

Module 6: Phi Data – Memory & Context Optimization

- **Introduction to Phi Data: AI Agent Memory & Learning**
- **Vector Database Integration (FAISS, Pinecone, ChromaDB)**
- **Retrieval-Augmented Generation (RAG) for AI Agents**
- **Personalized AI Assistants with Memory**

Hands-on:

- ✓ Implementing Phi Data for long-term memory
 - ✓ Storing and retrieving context for AI-powered Q&A
-

Module 7: Building & Deploying Agentic AI Applications

- **Deploying AI Agents on Cloud (Hugging Face, AWS, Azure)**
- **Integrating AI Agents into Web Applications (FastAPI, Streamlit)**
- **Security Considerations in Autonomous AI**



Mastering Prompt Engineering for LLMs

Module 1: Introduction to Prompt Engineering

- **What is Prompt Engineering?**
 - The role of prompts in LLMs
 - How LLMs process and interpret prompts
- **Why Learn Prompt Engineering?**
 - Optimizing LLM performance
 - Reducing hallucinations & improving accuracy
 - Enhancing AI-driven applications

Hands-on:

- ✓ Experimenting with OpenAI's API for basic text generation
-

Module 2: LLM Model Settings & Configurations

- **Understanding LLM Settings:**
 - Temperature, Top-k, Top-p (Nucleus Sampling)
 - Stop Tokens, Context Length & Tokenization
- **Fine-tuning vs Prompt Optimization**
- **Customizing Model Behavior with System Prompts**

Hands-on:

- ✓ Experimenting with different model settings in OpenAI Playground
-

Module 3: Prompt Elements & Structuring

- **Key Components of a Well-Designed Prompt:**

- Instructions
- Context
- Input Data
- Output Constraints
- **Optimizing Prompts for Accuracy & Consistency**

Hands-on:

- ✓ Structuring prompts for summarization, classification, and question-answering
-

Module 4: Prompt Engineering Techniques (Shot-Based Prompting)

- **Zero-Shot Prompting**
- **One-Shot Prompting**
- **Few-Shot Prompting**
- **Comparing Shot-Based Techniques for Different Use Cases**

Hands-on:

- ✓ Designing and testing different shot-based prompts on GPT models
-

Module 5: Chain of Thought (CoT) Prompting

- **What is CoT Prompting?**
- **Step-by-Step Reasoning in LLMs**
- **Implementing CoT in Math, Logic & Coding Tasks**

Hands-on:

- ✓ Using CoT to improve reasoning-based problem-solving
-

Module 6: Self-Consistency in Prompt Engineering

- **What is Self-Consistency?**
- **Generating Multiple Answers & Selecting the Best**
- **Improving Output Reliability with Self-Consistency**

Hands-on:

- ✓ Implementing Self-Consistency for multi-answer tasks
-

Module 7: Out-of-Date Learning in Prompt Engineering

- **How LLMs Handle Outdated Information**
- **Strategies to Overcome Out-of-Date Learning:**
 - Prompting with External Data
 - Fine-Tuning vs Retrieval-Augmented Generation (RAG)
- **When to Use Updated APIs & Tools**

Hands-on:

- ✓ Experimenting with model responses on time-sensitive queries
-

Module 8: Role-Playing in Prompt Engineering

- **What is Role-Playing in Prompting?**
- **Creating AI Personas for Specialized Tasks**
- **Enhancing Response Accuracy with Role-Based Prompts**


Hands-on:

- ✓ Designing AI assistants with different personas (e.g., Doctor, Lawyer, Coder)
-

Module 9: RAG (Retrieval-Augmented Generation) in Prompt Engineering

- What is RAG & Why is it Important?
- Integrating Vector Databases (FAISS, Pinecone, ChromaDB)
- Enhancing AI Responses with External Knowledge


Hands-on:

-  Implementing a RAG-based chatbot using FAISS & OpenAI
-

Module 10: ReAct (Reasoning + Acting) in Prompt Engineering

- What is ReAct Framework?
- Combining CoT + Tool Use for Autonomous Agents
- Building AI Agents that Reason & Execute Actions


Hands-on:

-  Implementing a ReAct-based agent using LangChain
-

Module 11: DSP (Dynamic Structured Prompting)

- What is Dynamic Structured Prompting (DSP)?
- Generating Structured & Dynamic Prompts Based on Context
- Using DSP for Adaptive AI Interactions

Hands-on:


-  Creating dynamically structured prompts for personalized AI responses

Advanced Course on Generative AI:

Module 1: Introduction to Generative AI

- What is Generative AI?
- Types of Generative AI Models:
 - Text-based (GPT, LLaMA, Claude)
 - Multimodal (CLIP, DALL-E, Stable Diffusion)
- Use Cases in NLP, Image Generation, and Code Generation


Hands-on:

-  Running a simple text-based generative model using OpenAI API
-

Module 2: Text-Based Generative Models

- How Text-Based Models Work
- Training LLMs (Large Language Models) with Transformers
- Pretrained Models vs Fine-Tuned Models
- Popular LLMs: GPT-4, LLaMA, Mistral, Falcon

Hands-on:

-  Generating text using Hugging Face Transformers
-

Module 3: Multimodal Models (Text + Image + Audio)

- What are Multimodal Models?
- Combining Text & Vision Models for AI Applications
- Examples of Multimodal Models: CLIP, DALL-E, Gemini, GPT-4 Turbo

Hands-on:

- ✓ Running OpenAI's CLIP model for text-to-image retrieval
-

📌 Module 4: CLIP (Contrastive Language-Image Pretraining) Architecture

- How CLIP Works: Text-Image Pairing
- Applications of CLIP in Image Search & Generation
- Fine-tuning CLIP for Custom Tasks

Hands-on:

- ✓ Using CLIP to find relevant images based on text prompts
-

📌 Module 5: VQGAN & Taming Transformers

- Introduction to VQGAN (Vector Quantized GAN)
- How Taming Transformers Improve Image Quality
- Combining VQGAN + CLIP for AI Art

Hands-on:

- ✓ Generating AI Art using VQGAN + CLIP
-

📌 Module 6: Autoencoders & VAEs (Variational Autoencoders)

- What is an Autoencoder?
- Difference Between Autoencoders & VAEs
- Generating High-Resolution Images with VAEs

Hands-on:

- ✓ Implementing a simple Variational Autoencoder (VAE) in PyTorch
-

📌 Module 7: Retrieval-Augmented Generation (RAG)

- What is RAG & Why It Matters for AI?
- Enhancing LLMs with External Knowledge
- Vector Databases for RAG (FAISS, Pinecone, ChromaDB)

Hands-on:

- ✓ Implementing a RAG-based chatbot using LlamaIndex & FAISS
-

📌 Module 8: Hugging Face Ecosystem

- Overview of Hugging Face Transformers
- Fine-tuning LLMs with Hugging Face
- Deploying Models Using Hugging Face Spaces

Hands-on:

- ✓ Fine-tuning a text generation model on Hugging Face
-

📌 Module 9: CrewAI for Multi-Agent AI Systems

- What is CrewAI?
- Building Teams of AI Agents
- Role-Based Task Assignment in CrewAI

Hands-on:

- ✓ Setting up an AI research team using CrewAI
-

📌 Module 10: Groq – High-Speed AI Inference

- What is Groq?

- Running AI Models at Lightning Speed
- Optimizing Large Models for Low Latency

Hands-on:

- ✓ Deploying a transformer model with Groq hardware
-

Module 11: Stable Diffusion for Image Generation

- Understanding Stable Diffusion Architecture
- Text-to-Image Generation with Diffusion Models
- Fine-Tuning & Customizing Stable Diffusion

Hands-on:

- ✓ Running Stable Diffusion on a local machine
-

Module 12: GitHub Copilot for AI-Powered Coding

- How GitHub Copilot Uses AI for Code Generation
- Best Practices for Using AI in Software Development
- Comparing Copilot with Other AI Coding Tools

Hands-on:

- ✓ Writing AI-assisted Python scripts using GitHub Copilot
-

Module 13: LlamaIndex – AI-Powered Document Processing

- What is LlamaIndex?
- Connecting LLMs to Private Data
- Using LlamaIndex for Enterprise AI Applications

Hands-on:

- ✓ Implementing LlamaIndex for a document-based AI assistant
-

Module 14: FastAPI for AI Model Deployment

- Introduction to FastAPI for AI
- Building a REST API for LLMs
- Deploying AI Models as Web Services

Hands-on:

- ✓ Deploying a text-based LLM using FastAPI

Advanced Course on : LLMs

Module 1: Introduction to Generative AI & LLMs

- What is Generative AI?

- Types of Generative AI Models (Text, Image, Multimodal, Speech)
- Comparison of Leading AI Models (GPT, Gemini, LLaMA, Claude, Mixtral, DeepSeek, Grok)
- Foundation Models vs Fine-Tuned Models

Hands-on:


- ✓ Running a basic LLM-powered chatbot using OpenAI API
-

Module 2: OpenAI's AI Ecosystem

LangChain Framework

- Introduction to LangChain
- Building AI Agents with LangChain
- Memory & Context Handling in LangChain
- Connecting LLMs with External Data Sources

Hands-on:

- ✓ Implementing a chatbot using OpenAI + LangChain
-  OpenAI Whisper (Speech-to-Text AI)
 - What is OpenAI Whisper?
 - Multilingual Speech Recognition
 - Building Real-World Applications with Whisper

Hands-on:


- ✓ Transcribing audio into text using OpenAI Whisper
-

Module 3: Google's Generative AI Ecosystem

Gemini AI: Google's Multimodal LLM

- Introduction to Gemini AI
- Comparison with OpenAI & Meta Models
- Using Gemini API for Text & Image Generation

Hands-on:

- ✓ Generating text & images using Gemini API
-  Google Vision: AI for Image Analysis & Recognition
 - What is Google Vision?
 - AI-Powered Image Processing & OCR
 - Building AI-powered Image Search Systems

Hands-on:


- ✓ Using Google Vision API for image classification
-

Module 4: DeepSeek, Mistral, Mixtral, Grok & Claude

DeepSeek AI

- Overview of DeepSeek Language Model
- Optimizing Search and AI Retrieval with DeepSeek
- Applications in AI-Assisted Knowledge Systems

Hands-on:

- ✓ Running DeepSeek for document-based AI search
-  Mistral & Mixtral: High-Performance Open-Source AI
 - What is Mistral & Mixtral?
 - Dense vs Sparse Transformer Models
 - Optimizing Mixtral for Multi-Task AI

Hands-on:

- ✓ Running Mixtral models on Hugging Face
- Grok (X AI by Elon Musk)
 - How Grok is Designed for Real-Time AI Processing
 - Comparison with GPT, Gemini & Claude
 - Use Cases of Grok in AI Chatbots & Assistants

Hands-on:

- ✓ Running Grok on X (formerly Twitter) API
- Claude (Anthropic AI)
 - What is Claude & How It Differs from Other LLMs?
 - Claude's Approach to AI Safety & Constitutional AI
 - Fine-Tuning Claude for Enterprise Applications

Hands-on:

- ✓ Building an AI-powered assistant with Claude API
-

📌 Module 5: Meta's Generative AI Ecosystem

- LLaMA 3: Open-Source LLM by Meta
 - Introduction to LLaMA 3
 - Comparison with GPT-4, Gemini, Mixtral, & Claude
 - Fine-Tuning & Customizing LLaMA for Specific Tasks

Hands-on:

- ✓ Running LLaMA 3 on a local machine using Hugging Face
- Building Generative AI on Cloud
 - Cloud Platforms for AI (AWS, GCP, Azure, Meta Cloud)
 - Deploying LLMs on Cloud for Scalability
 - Building AI-Driven Web Apps with Cloud-Based LLMs

Hands-on:

- ✓ Deploying an LLM-powered chatbot on Cloud
- Meta's Foundation Models
 - Understanding Meta's AI Foundation Models
 - Pre-Trained Models vs Custom Models
 - Adapting Foundation Models for Industry Use Cases

Hands-on:

- ✓ Using Meta's AI Models for custom NLP tasks
-

📌 Module 6: Fine-Tuning LLMs with Quantization, LoRA & QLoRA

- ◆ Fine-Tuning Large Language Models (LLMs)
 - Why Fine-Tune an LLM?
 - Datasets & Preprocessing for LLM Fine-Tuning
 - Fine-Tuning vs Prompt Engineering
- ◆ LoRA (Low-Rank Adaptation) Fine-Tuning
 - What is LoRA?
 - Reducing Computation for LLM Training
 - Implementing LoRA with Hugging Face & PyTorch
- ◆ QLoRA (Quantized LoRA) for Efficient Model Fine-Tuning
 - What is QLoRA?

- Memory Optimization for Large LLMs
- Running Fine-Tuned Models on Low-End Hardware

Hands-on:

- ✓ Fine-tuning a LLaMA 3 model using LoRA & QLoRA

📌 Module 7: End-to-End AI Model Deployment & Optimization

- Deploying AI Models with APIs (FastAPI, Flask)
- Optimizing AI Models for Performance & Cost
- Best Practices for AI Model Security & Governance

Hands-on:

- ✓ Deploying a fine-tuned LLM as a FastAPI web service

Vector database

📌 Module 1: Introduction to Vector Databases

What is a Vector Database?

Difference Between Traditional & Vector Databases

Why Use Vector Databases in Generative AI & LLMs?

How Vector Embeddings Work in AI Search & Retrieval

Hands-on:

- ✓ Generating vector embeddings from text using OpenAI's text-embedding-ada-002

📌 Module 2: Understanding Vector Embeddings

What Are Embeddings in AI?

How LLMs Convert Text, Images & Audio to Vectors

Similarity Metrics: Cosine Similarity, Euclidean Distance, Dot Product

Choosing the Right Embedding Model (OpenAI, Hugging Face, SentenceTransformers, BERT, etc.)

Hands-on:

- ✓ Generating embeddings with OpenAI, Hugging Face, and BERT models

📌 Module 3: Implementing FAISS (Facebook AI Similarity Search)

What is FAISS & How It Works?

Indexing & Searching Large-Scale Vectors with FAISS

Optimizing FAISS for Fast Retrieval

Hands-on:

- ✓ Implementing a FAISS-based search engine for document retrieval

📌 Module 4: Using Pinecone for Scalable AI Search

Introduction to Pinecone: A Managed Vector Database

Building Real-Time AI Search Applications with Pinecone

Comparing FAISS vs Pinecone vs Milvus

Hands-on:

- ✓ Creating a question-answering chatbot using OpenAI + Pinecone

📌 Module 5: ChromaDB for LLMs & AI Applications

What is ChromaDB?

How ChromaDB Works with LangChain

Building RAG (Retrieval-Augmented Generation) Pipelines with ChromaDB

Hands-on:

- ✓ Integrating ChromaDB with OpenAI's GPT for AI-powered search

📌 Module 6: Exploring Weaviate & Its AI Capabilities

Overview of Weaviate as a Hybrid Search Engine

Using Weaviate for Semantic Search & Knowledge Graphs

Deploying Weaviate on Cloud & Local Environments

Hands-on:

- ✓ Implementing a semantic search engine with Weaviate

📌 Module 7: Milvus & Qdrant for Large-Scale AI Applications

● Milvus (Distributed & Cloud-Based Vector Search)

Introduction to Milvus for AI Applications

Optimizing Milvus for Large-Scale Data Processing

● Qdrant (High-Performance Open-Source Vector DB)

What is Qdrant & How It Works?

Fine-Tuning Qdrant for AI Search & Recommendation Systems

Hands-on:

- ✓ Deploying Milvus & Qdrant for AI-driven search and recommendations

📌 **Module 8: Building AI-Powered Search & RAG Applications**
What is Retrieval-Augmented Generation (RAG)?

Integrating Vector Databases with LLMs for Intelligent Search

Building Enterprise AI Assistants Using Vector Databases

Hands-on:

✅ Building an RAG pipeline using LangChain, Pinecone, and OpenAI

MLOps & CI/CD Pipeline for AI & Machine Learning

📌 **Module 1: Introduction to MLOps & CI/CD for ML Models**

- What is MLOps & Why is it Important?
- DevOps vs MLOps: Key Differences
- MLOps Lifecycle & Stages
- Understanding CI/CD Pipelines in AI/ML Projects
- Challenges in Deploying ML Models at Scale

Hands-on:

✅ Setting up a basic **CI/CD pipeline for ML models** using GitHub Actions

📌 **Module 2: Containerization & Model Packaging with Docker**

- Why Containerize ML Models?
- Building & Running Docker Containers for ML Applications
- Deploying ML Models inside Docker Containers
- Optimizing Containers for AI Workloads

Hands-on:

✅ Containerizing an ML Model with Docker & Running it Locally

📌 **Module 3: CI/CD Pipeline for ML with GitHub Actions & Jenkins**

- What is CI/CD & How Does It Work in ML?
- Setting up a CI/CD Pipeline for ML Models
- Automating Model Testing, Validation & Deployment
- CI/CD with GitHub Actions,

Hands-on:

✅ Implementing a CI/CD Pipeline for an AI Model Deployment Using GitHub Actions

📌 **Module 4: MLOps with MLflow**

- Introduction to Kubeflow for AI & ML
- Integrating with MLflow for Experiment Tracking

Hands-on:

 **Building an ML Workflow with Kubeflow Pipelines**
